

Case Studies

Wednesday, September 9, 2020 11:52 AM

① Why case studies

→ Classic Networks

- LeNet - 5
- AlexNet
- VGG

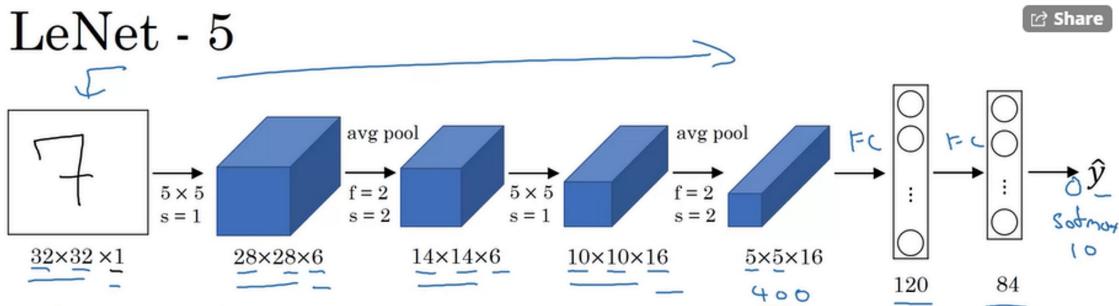
→ ResNet Network

→ Inception Network

② Classic Networks

→ LeNet - 5

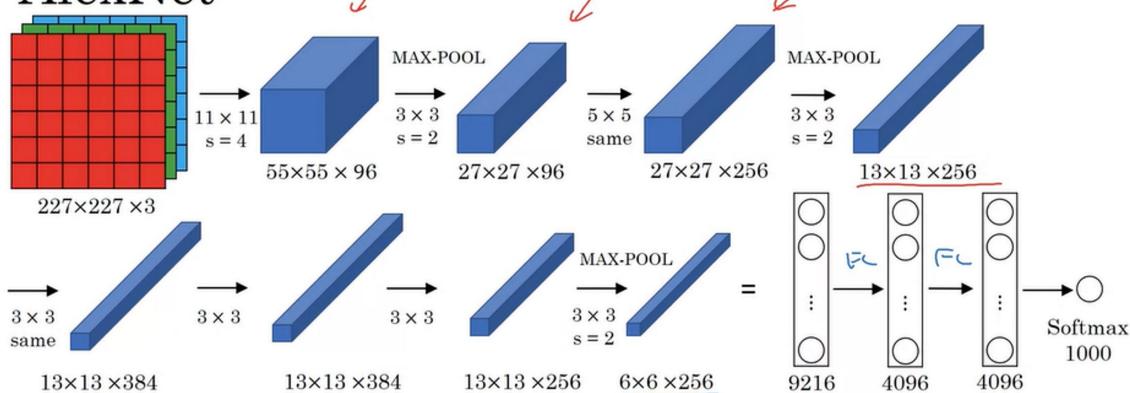
- To recognise handwritten digits
- $32 \times 32 \times 1$ images (a grayscale image)
- Padding wasn't used
- Small for today's standards (~60K parameters)
- $n_k \downarrow$ & $n_c \uparrow$ as we go deeper
- conv + pool → conv + pool → fc → fc → output



→ AlexNet

- Similar to LeNet, but much bigger
- 60M parameters
- used ReLU activation function
- multiple GPUs were used in training

AlexNet



→ VGG - 16

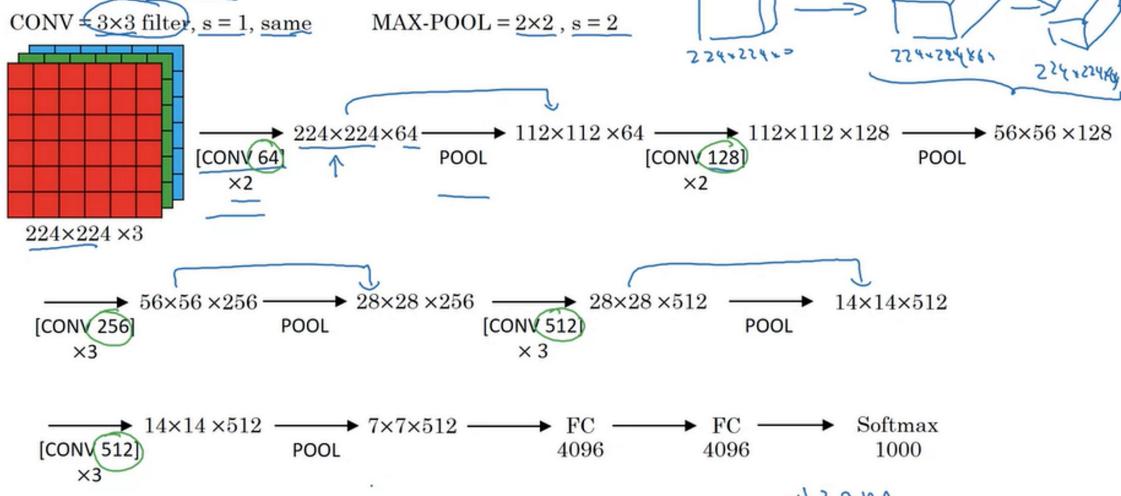
conv + pool + fc

13x13x384 13x13x384 13x13x256 6x6x256 9216 4096 4096

\rightarrow VGG-16

- Deeper than previous networks
- ~138M parameters
- Filters were doubled after every pool
- The uniformity of this network made it attractive

VGG - 16



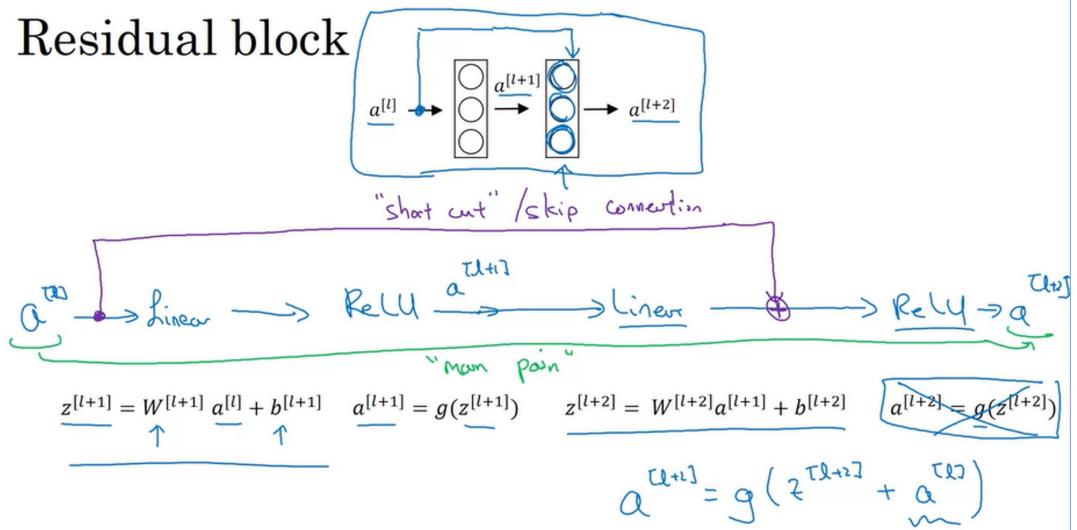
[Simonyan & Zisserman 2015. Very deep convolutional networks for large-scale image recognition]

Andrew Ng

③ ResNets

- Using skip connections

Residual block

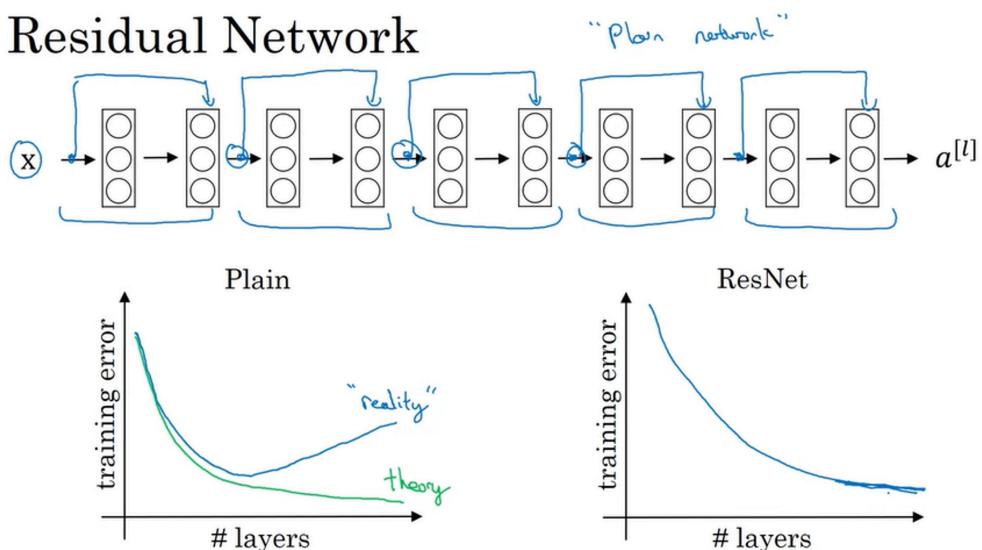


[He et al., 2015. Deep residual networks for image recognition]

Andrew Ng

- Residual Network

Residual Network



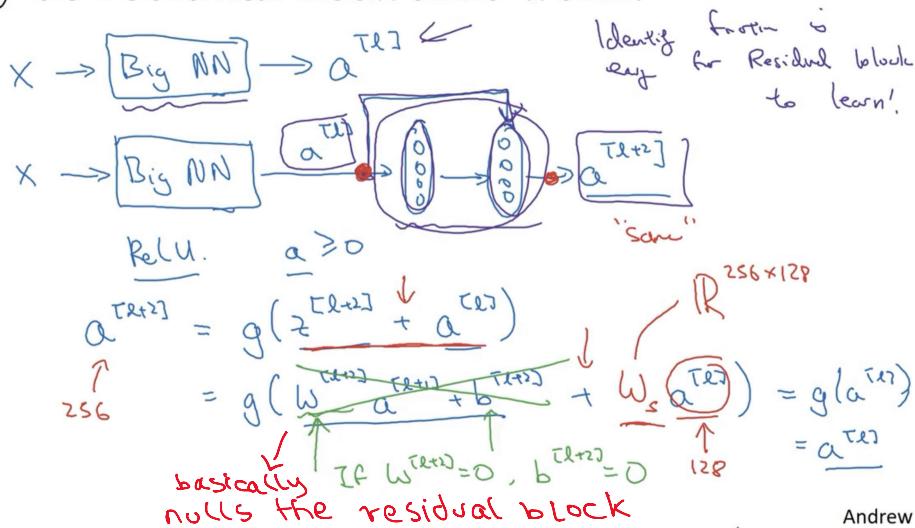
[He et al., 2015. Deep residual networks for image recognition]

Andrew Ng

④ Why ResNets work

→ Usually deeper networks hurt performance
But this isn't true for ResNets

Why do residual networks work?

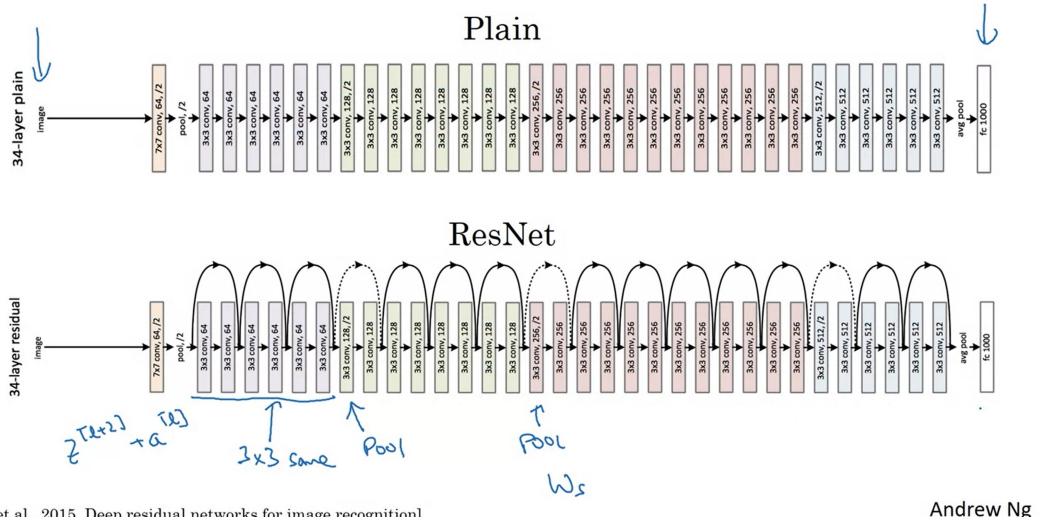


Andrew Ng

- Using a residual block never hurts performance as it can easily learn the identity fn.
- This is particularly helpful in large networks where lot of parameters make it hard for the network to learn identity function

→ ResNet Paper

ResNet



⑤ Networks in Networks and 1×1 Convolutions

→ When using 1×1 convolution,

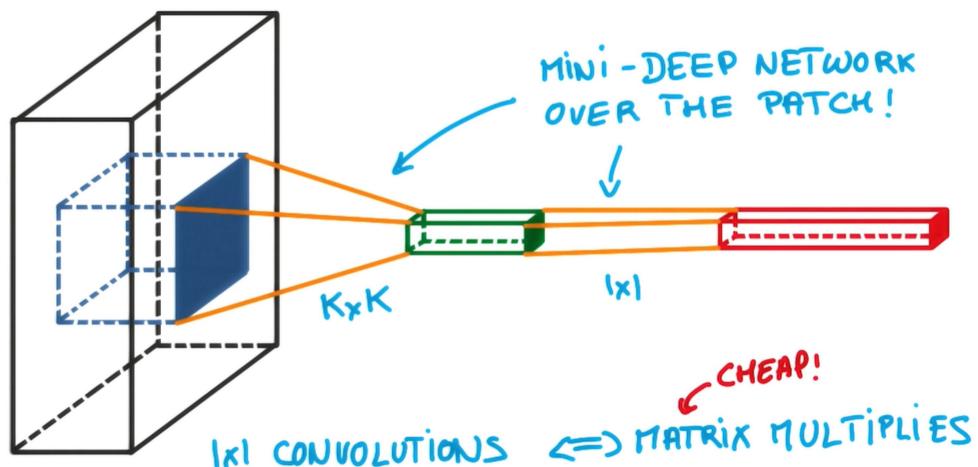
- Input: $6 \times 6 \times 1$
Conv: $1 \times 1 \times 1$ (one filter)
Output: $6 \times 6 \times 1$
- Input: $6 \times 6 \times 32$
Conv: $1 \times 1 \times 32$ (5 filters)
Output: $6 \times 6 \times 5$

→ It is useful when

- We want to shrink number of channels called feature transformation
- This can help in saving lot of computations
- If we use same number of filters as number of channels, output contains same number of channels.
- The 1×1 conv will act like a non-linearity and will learn non-linearity operator

→ We could replace a fully connected layer with 1×1 conv as they are essentially the same

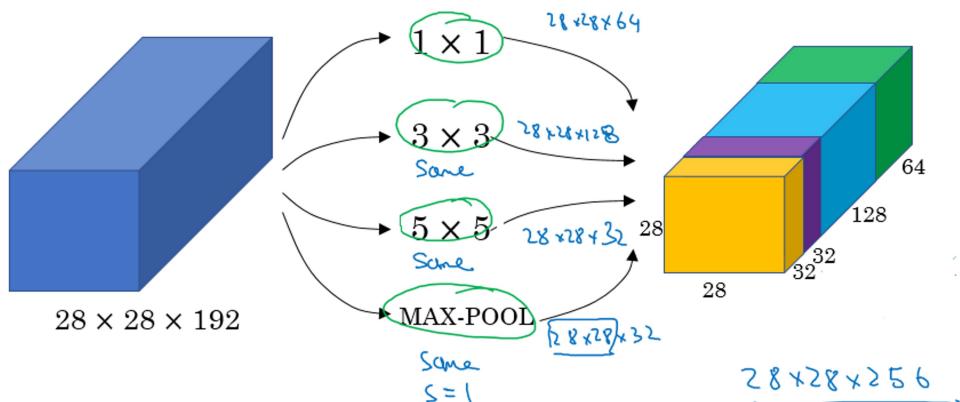
1x1 CONVOLUTIONS



→ Cheap way to make model deeper without adding lot of parameters

⑥ Inception Network Motivation

→ Use multiple convolutions and let model decide weights



→ Problem of computational cost

- Since we use multiple filters the no. of operations is very high
consider just 5x5 conv.
we use 32 filters

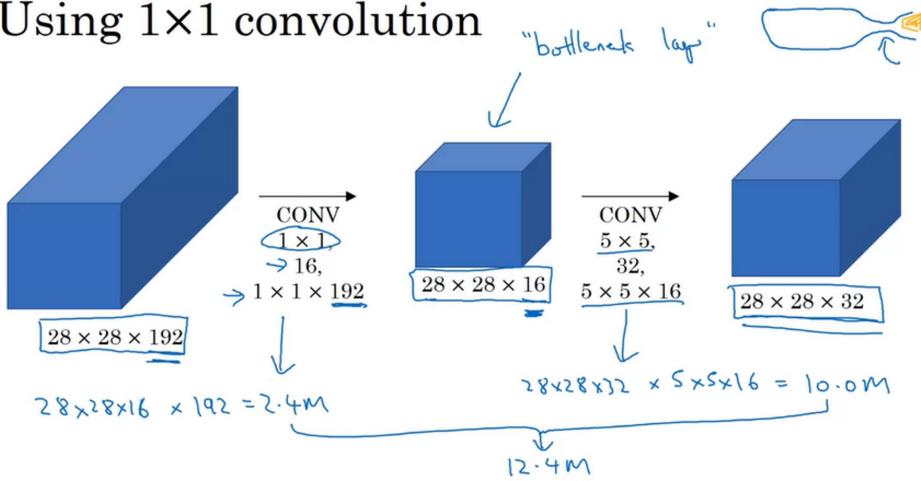
$$\begin{aligned}\therefore \text{No. of operations} &= \text{no. of outputs} \times F \times F \times \text{Input dimensions} \\ &= (28 \times 28 \times 32) \times (5 \times 5) \times 192 \\ &= \text{c20 Million}\end{aligned}$$

- 120 Million is large even for modern computers
- Using 1x1 conv.
Using 1x1 convolution

"bottleneck layer"



Using 1×1 convolution



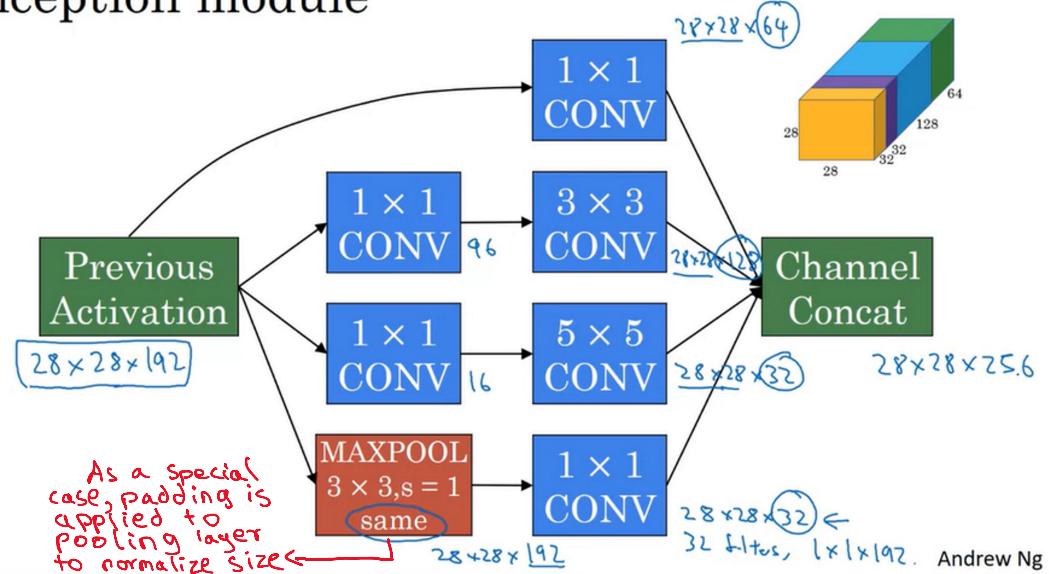
we have reduced parameters from $10.0M$ to $2.4M$ by using a 1×1 "bottleneck layer"

→ shrinking the representation size doesn't affect performance of model as long as done within reason

⑦ Inception Network

→ Inception Module

Inception module



→ The inception network combines a lot of these modules to make a deep network

GoogleNet

